

10 MISTAKES

10 AI Agent Mistakes

That Kill Your Interview in 5 Minutes

Interviewers test systems, not prompts. These mistakes end interviews in minutes.

Agents

Architecture

Security

Production

Interviews



Naved Khan
Senior Gen-AI Engineer

+ Follow

Developers often put planning, memory, and tool usage into one agent. That lack of modularity makes the system brittle and impossible to debug.

- ▶ **Bottlenecks:** Single failure crashes the whole flow
- ▶ **Latency:** Huge prompts slow down even simple tasks
- ▶ **Debugging:** You can't trace which part failed
- ▶ **Scaling:** You can't optimize one component (e.g. a search tool) when everything is coupled

WHY IT MATTERS

Split into smaller, single-purpose agents so you can test, fix, and scale each part on its own.



Naved Khan
Senior Gen-AI Engineer

+ Follow

You secure user input but forget that agents read external data. Malicious instructions hidden on a webpage or in a doc can hijack the agent.

- ▶ **Data exfiltration:** Agent sends logs or data to an attacker
- ▶ **Privilege escalation:** Agent uses tools it shouldn't
- ▶ **Integrity loss:** Agent deletes or changes records from web instructions

DEFENSE

Add firewalls at the agent-tool boundary: sanitize or minimize data from tools before the agent sees it. Don't trust raw external content.



Naved Khan
Senior Gen-AI Engineer

[+ Follow](#)

In multi-agent systems, one agent's small hallucination becomes "ground truth" for the next. Errors compound at machine speed.

- ▶ **Compounded hallucinations:** By step 10 the output is fully wrong
- ▶ **Invisible failures:** Output looks plausible but is flawed
- ▶ **Audit:** Hard to find which agent introduced the error
- ▶ **Waste:** Tokens and compute spent on dead-end paths from early errors

WHY IT MATTERS

Add validation or human checkpoints between critical steps. Log intermediate outputs so you can trace and fix the source.



Naved Khan
Senior Gen-AI Engineer

[+ Follow](#)

Relying on "infinite context" instead of structured memory leads to poisoning. Old, irrelevant data starts driving decisions.

- ▶ **Circular reasoning:** Agent repeats past mistakes from context
- ▶ **Cost:** Context-heavy prompts cost more per call
- ▶ **Contradictions:** Agent follows outdated instructions still in context
- ▶ **UX:** Agent "forgets" real user preferences but keeps useless noise

WHY IT MATTERS

Use structured memory: separate short-term context from long-term facts. Summarize and evict instead of dumping everything into the prompt.



Naved Khan
Senior Gen-AI Engineer

+ Follow

Giving an agent broad API access "to save time" is a huge risk. A CRM agent with admin rights can cause serious damage if it misbehaves.

- ▶ **Unauthorized access:** Surfacing HR or sensitive data
- ▶ **Irreversible actions:** Deleting production DBs or records
- ▶ **Compliance:** Breaking GDPR or other rules
- ▶ **Attack surface:** Every tool the agent uses is a potential exploit path

WHY IT MATTERS

Grant the least privilege. Each tool should have the minimum scope needed. Validate permissions inside the tool layer, not only in prompts.



Naved Khan
Senior Gen-AI Engineer

+ Follow

The same prompt can yield different results. That breaks downstream systems that expect a fixed JSON schema or format.

- ▶ **System crashes:** Wrong format breaks the UI or parser
- ▶ **Inconsistent CX:** Same question gets different answers
- ▶ **Fragility:** Small phrase changes make parsers fail
- ▶ **Testing:** Unit tests don't catch these "soft" failures; you need eval frameworks

WHY IT MATTERS

Use structured outputs (JSON mode, tool schemas) and validation. Add evaluation runs that check format and consistency, not only correctness.



Naved Khan
Senior Gen-AI Engineer

+ Follow

Autopilot for financial or legal actions is dangerous. Sensitive calls need checkpoint UIs where a human approves before the action runs.

- ▶ **Financial loss:** Wrong trades, refunds, or transfers
- ▶ **Legal liability:** Agreeing to bad terms or commitments
- ▶ **Reputation:** Offensive or wrong public actions
- ▶ **Safety:** In IoT or robotics, lack of oversight can cause physical harm

WHY IT MATTERS

Design approval flows for high-impact actions. Never let the agent execute irreversible or high-stakes steps without a human gate.



Naved Khan
Senior Gen-AI Engineer

+ Follow

Logging only the final output hides failures. You must log intermediate steps: which tools were called, with what inputs, and what they returned.

- ▶ **Silent failures:** You can't tell if a tool failed or was skipped
- ▶ **ROI:** You can't see which tools add value
- ▶ **Compliance:** Reasoning is opaque for audits
- ▶ **Optimization:** You can't find slow or error-prone tools to fix

WHY IT MATTERS

Log tool calls, inputs, and outputs. Use traces for debugging and for evaluating which tools are reliable and worth keeping.



Naved Khan
Senior Gen-AI Engineer

+ Follow

LLMs and APIs fail. Without graceful degradation, an API timeout or hallucination spike leaves the user with a dead app.

- ▶ **Total downtime:** One outage kills the whole product
- ▶ **Infinite loops:** Agent burns budget in minutes retrying
- ▶ **Poor UX:** "Server Error" instead of a helpful message or fallback path
- ▶ **Gridlock:** One stuck agent can block other system processes

WHY IT MATTERS

Add timeouts, retry limits, and fallback responses. When the agent or a tool fails, show a clear message and optional manual path.



Naved Khan
Senior Gen-AI Engineer

+ Follow

Prompts tuned for one model version can break on the next. Accuracy can drop overnight without monitoring.

- ▶ **Logic breakage:** Accuracy drops suddenly after a model update
- ▶ **Obsolescence:** Complex prompts stop working as intended
- ▶ **Tool errors:** Model stops following conventions or schemas
- ▶ **Maintenance:** Teams re-tune the whole system every few months

WHY IT MATTERS

Monitor accuracy and behavior on a sample of real traffic. Pin model versions where possible and test new versions before rollout.



Naved Khan
Senior Gen-AI Engineer

+ Follow

Want more content like this?



Naved Khan

Senior Gen-AI Engineer

Follow for more AI and Gen-AI content



Like



Comment



Share



Save